

Assignment -3
Problem Statement-Abalone Age Prediction

| | |
|---------------------|-----------------|
| Assignment Date | 06 October 2022 |
| Student Name | Miss.Gowshika N |
| Student Roll Number | 620119106024 |
| Maximum Marks | 2 Marks |

Question-1:

1.download the data set

2.load the dataset

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn
```

```
data=pd.read_csv("abalone.csv")data.head()
```

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0 | M | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 15 |
| 1 | M | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 7 |
| 2 | F | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 9 |
| 3 | M | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 10 |
| 4 | I | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 7 |

```
Age=1.5+data.Ringsdata["Age"]=Agedata=data.rename(columns = {'Whole weight':'Whole_weight','Shucked weight': 'Shucked_weight','Viscera weight': 'Viscera_weight',
'Shell weight': 'Shell_weight'})data=data.drop(columns=["Rings"],axis=1)data.head()
```

| | Sex | Length | Diameter | Height | Whole_weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|------|
| 0 | M | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 16.5 |
| 1 | M | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 8.5 |
| 2 | F | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 10.5 |
| 3 | M | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 11.5 |
| 4 | I | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 8.5 |

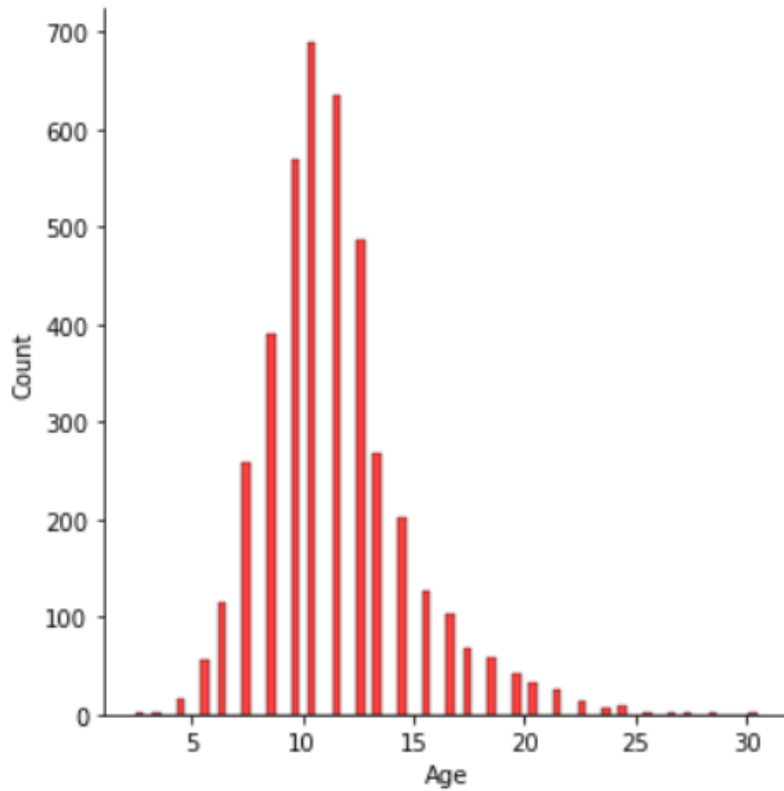
3.perform below visualizations
univariate analysis

Ans:

```
sns.displot(data["Age"], color='red')
```

Out[4]:

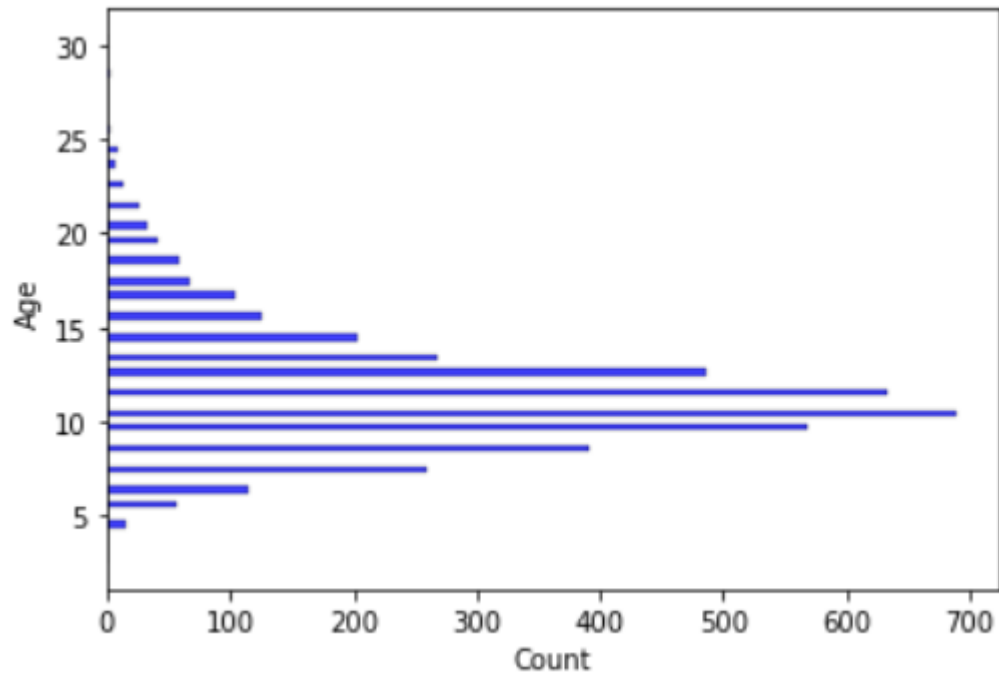
<seaborn.axisgrid.FacetGrid at 0x187953d26a0>



```
sns.histplot(y=data.Age,color='blue')
```

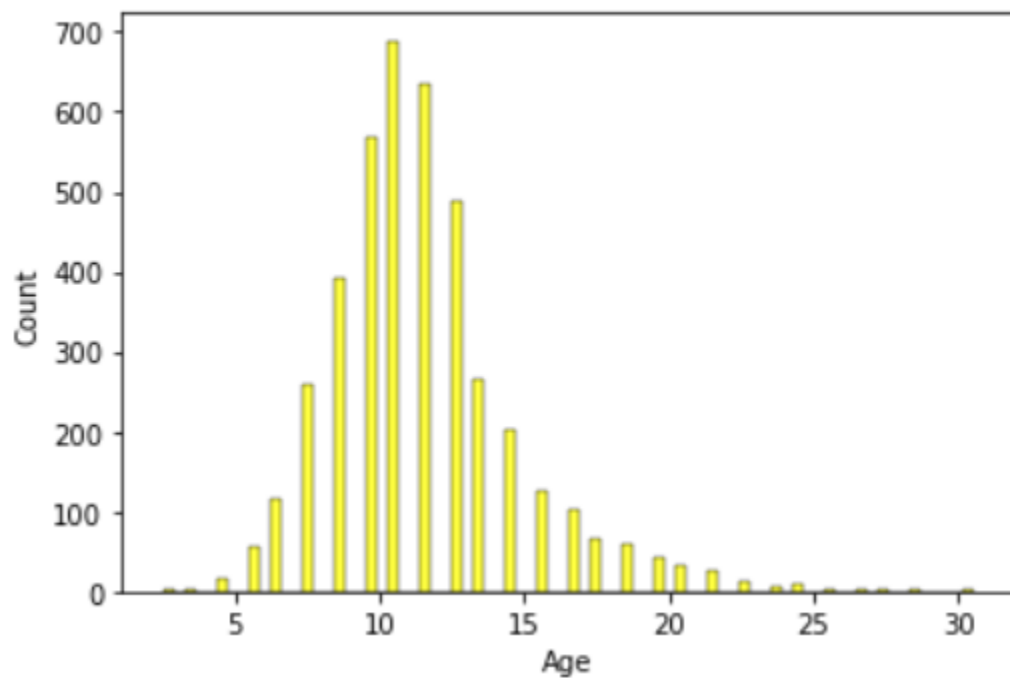
Out[5]:

<AxesSubplot:xlabel='Count', ylabel='Age'>



```
sns.histplot(x=data.Age,color='yellow')
```

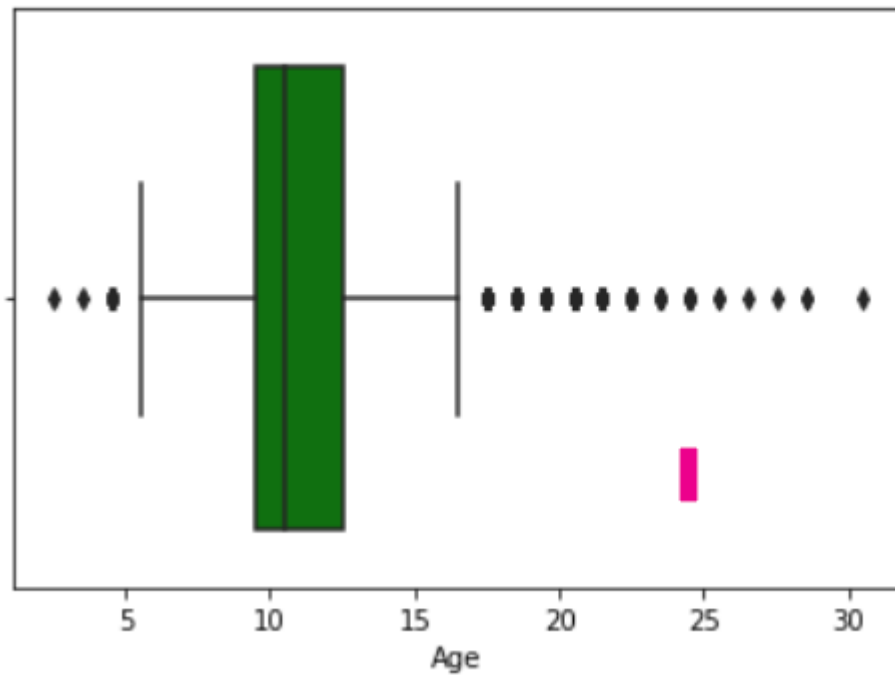
```
<AxesSubplot:xlabel= 'Age', ylabel= 'Count'>
```



In [6]:

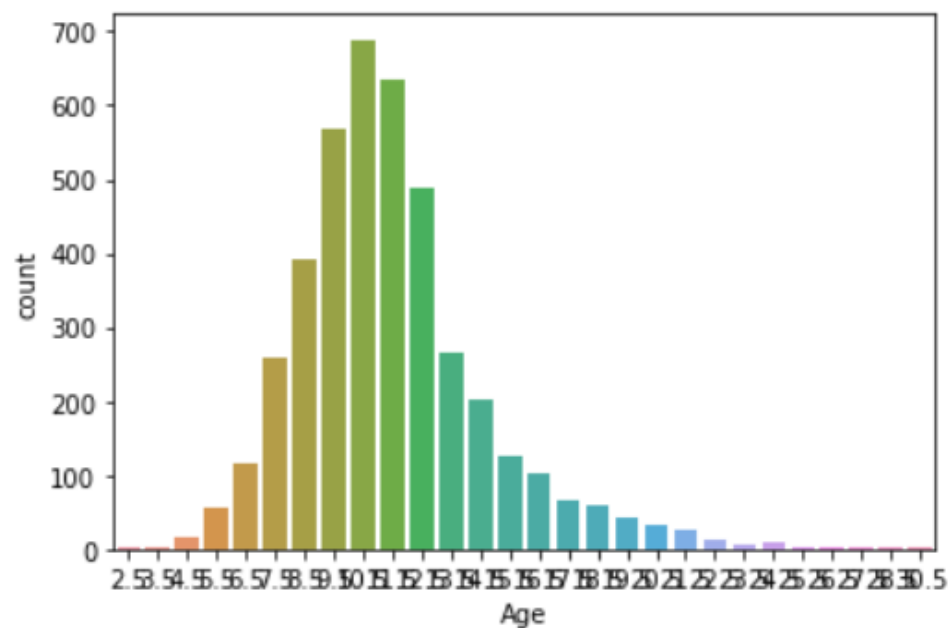
```
sns.boxplot(x=data.Age,color='green')
```

```
<AxesSubplot:xlabel='Age'>
```



```
sns.countplot(x=data.Age)
```

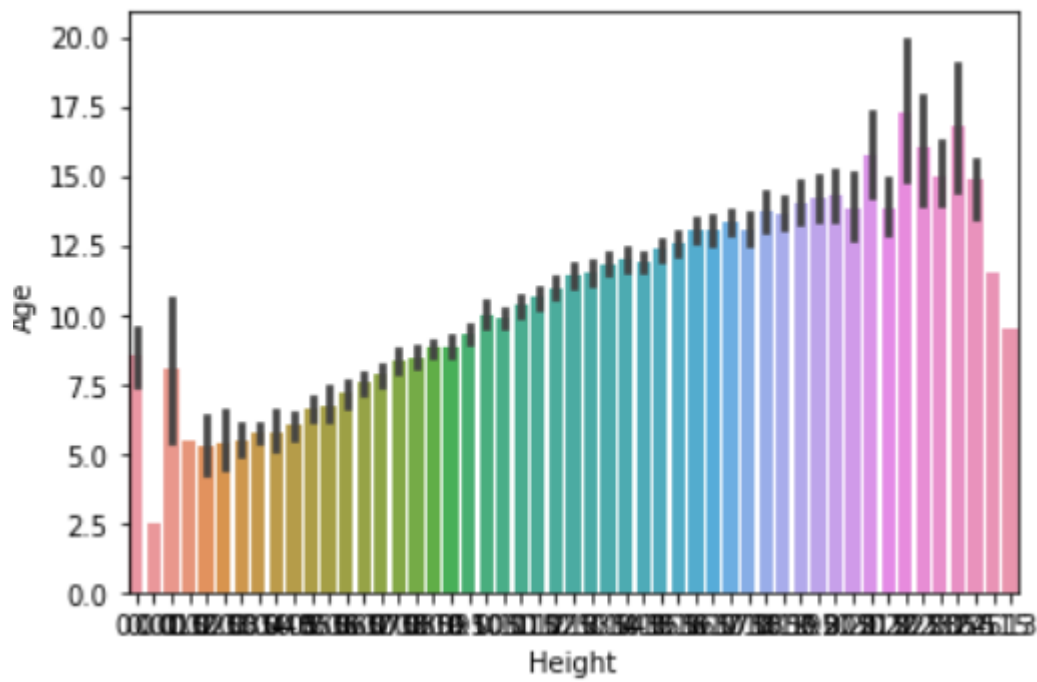
```
<AxesSubplot:xlabel='Age', ylabel='count'>
```



bi - variate analysis

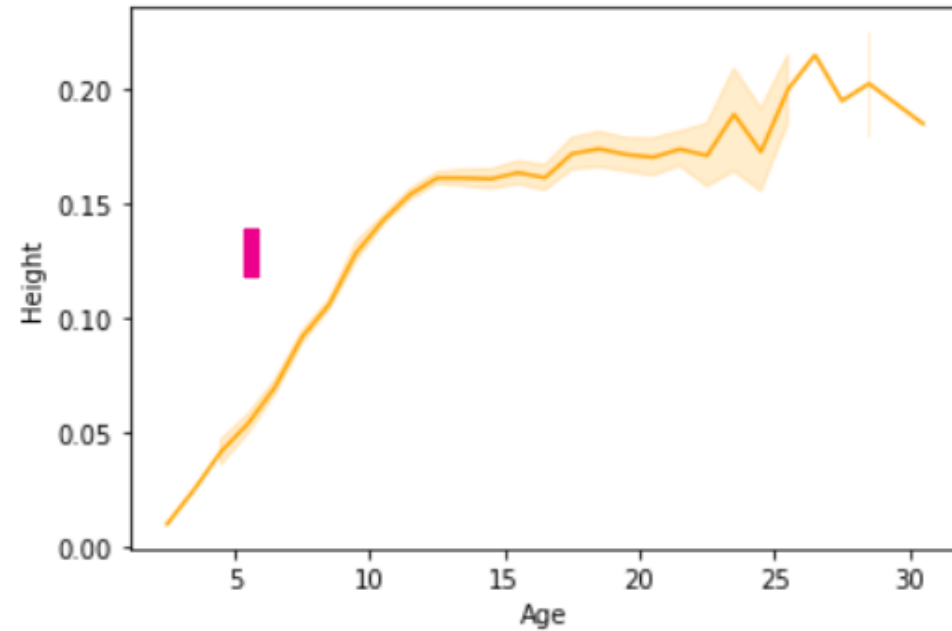
```
sns.barplot(x=data.Height,y=data.Age)
```

```
<AxesSubplot:xlabel='Height', ylabel='Age'>
```



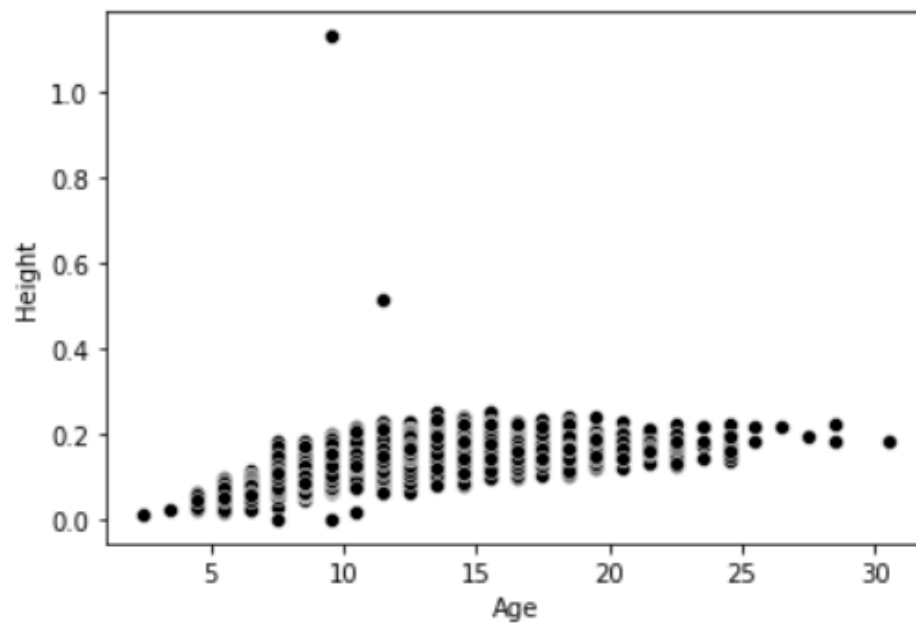
```
sns.lineplot(x=data.Age,y=data.Height, color='orange')
```

```
<AxesSubplot:xlabel='Age', ylabel='Height'>
```



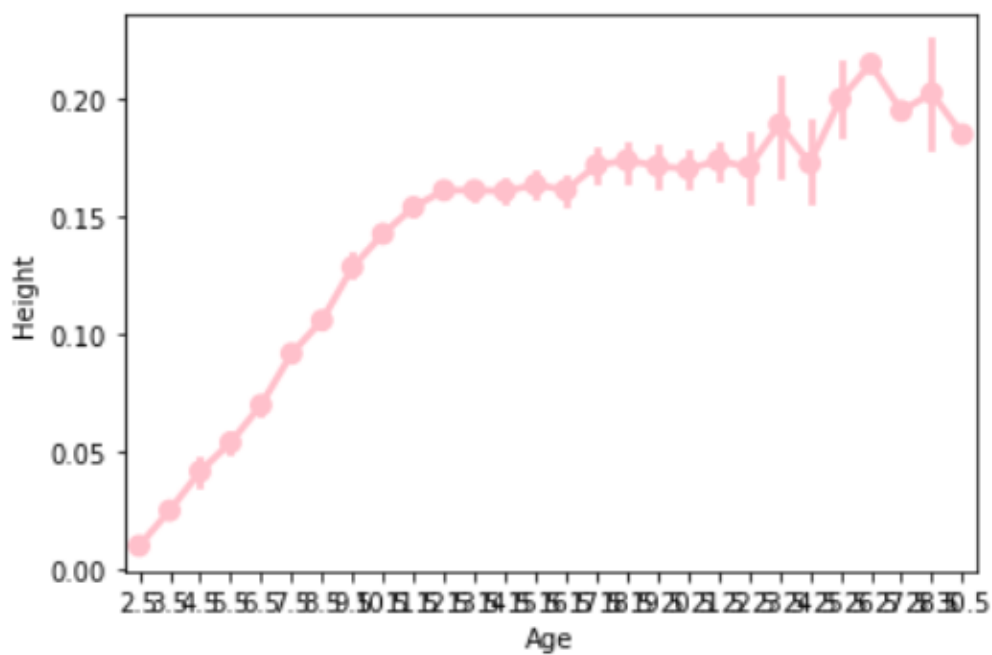
```
sns.scatterplot(x=data.Age,y=data.Height,color='black')
```

```
<AxesSubplot:xlabel='Age', ylabel='Height'>
```



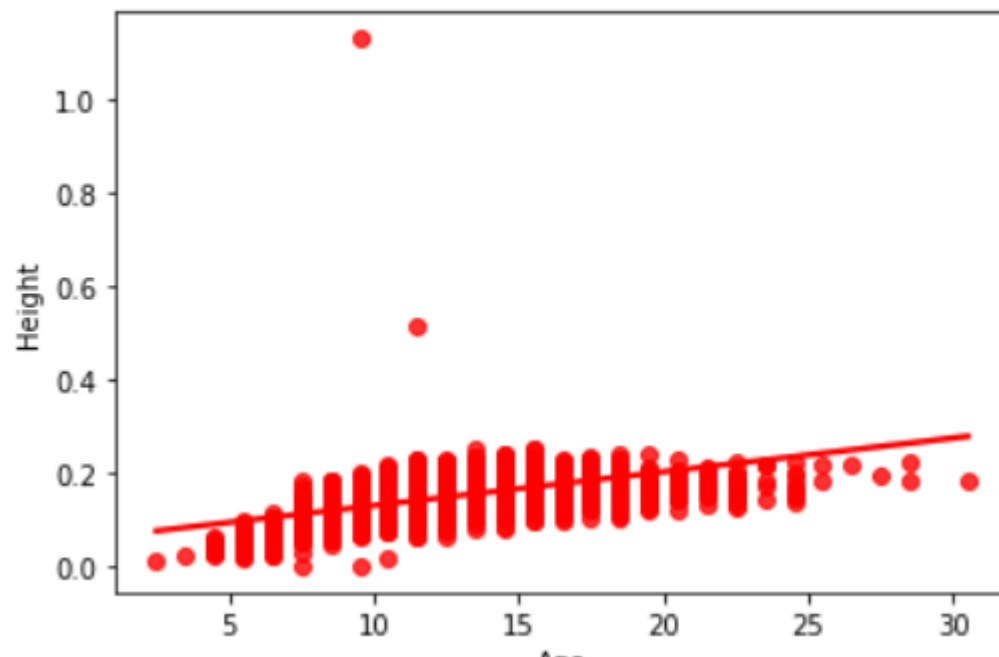
```
sns.pointplot(x=data.Age, y=data.Height, color="pink")
```

```
<AxesSubplot:xlabel='Age', ylabel='Height'>
```



```
sns.regplot(x=data.Age, y=data.Height, color='red')
```

```
<AxesSubplot:xlabel='Age', ylabel='Height'>
```

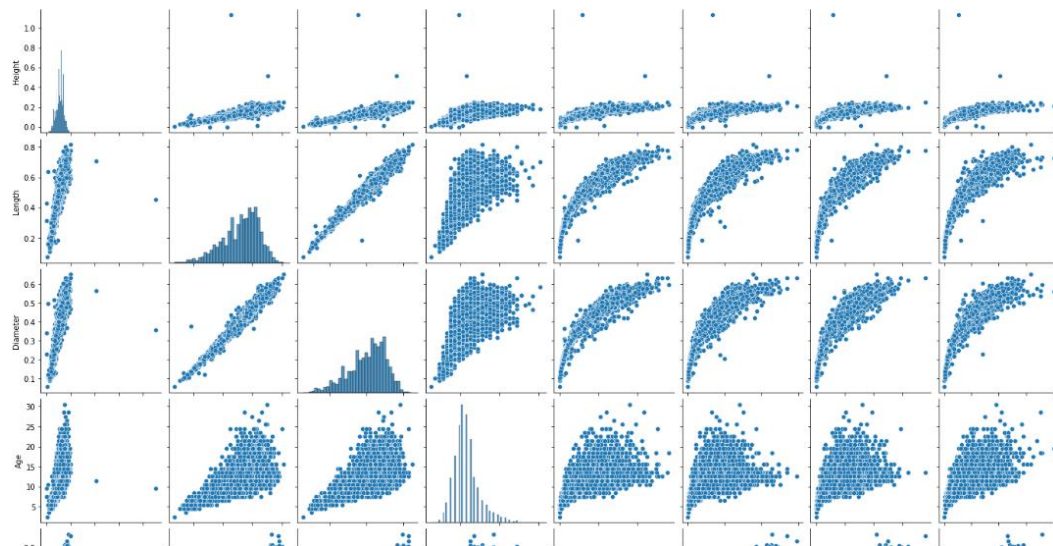


multi - variate analysis

In [9]:

```
sns.pairplot(data=data[["Height","Length","Diameter","Age","Whole_weight","Shucked_weight","Viscera_weight","Shell_weig
```

```
15]: <seaborn.axisgrid.PairGrid at 0x18795288400>
```



4. descriptive statistics

```
data.describe(include='all')
```

| | Sex | Length | Diameter | Height | Whole_weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|--------|------|-------------|-------------|-------------|--------------|----------------|----------------|--------------|-------------|
| count | 4177 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 |
| unique | 3 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| top | M | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| freq | 1528 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| mean | NaN | 0.523992 | 0.407881 | 0.139516 | 0.828742 | 0.359367 | 0.180594 | 0.238831 | 11.433684 |
| std | NaN | 0.120093 | 0.099240 | 0.041827 | 0.490389 | 0.221963 | 0.109614 | 0.139203 | 3.224169 |
| min | NaN | 0.075000 | 0.055000 | 0.000000 | 0.002000 | 0.001000 | 0.000500 | 0.001500 | 2.500000 |
| 25% | NaN | 0.450000 | 0.350000 | 0.115000 | 0.441500 | 0.186000 | 0.093500 | 0.130000 | 9.500000 |
| 50% | NaN | 0.545000 | 0.425000 | 0.140000 | 0.799500 | 0.336000 | 0.171000 | 0.234000 | 10.500000 |
| 75% | NaN | 0.615000 | 0.480000 | 0.165000 | 1.153000 | 0.502000 | 0.253000 | 0.329000 | 12.500000 |
| max | NaN | 0.815000 | 0.650000 | 1.130000 | 2.825500 | 1.488000 | 0.760000 | 1.005000 | 30.500000 |

5. Check for Missing values and deal with them

```
data.isnull().sum()
```

```
Sex          0
Length       0
Diameter     0
Height       0
Whole_weight 0
Shucked_weight 0
Viscera_weight 0
Shell_weight 0
Age          0
dtype: int64
```

6. Find the outliers and replace them

outliers

```
In [18]:
```


```
outliers=data.quantile(q=(0.25,0.75))outliers
```


| | Length | Diameter | Height | Whole_weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|-------------|--------|----------|--------|--------------|----------------|----------------|--------------|------|
| 0.25 | 0.450 | 0.35 | 0.115 | 0.4415 | 0.186 | 0.0935 | 0.130 | 9.5 |
| 0.75 | 0.615 | 0.48 | 0.165 | 1.1530 | 0.502 | 0.2530 | 0.329 | 12.5 |

```
a = data.Age.quantile(0.25)b = data.Age.quantile(0.75)c = b - a
lower_limit = a - 1.5 * c
data.median(numeric_only=True)
```

```
Length          0.5450
Diameter        0.4250
Height          0.1400
Whole_weight    0.7995
Shucked_weight  0.3360
Viscera_weight  0.1710
Shell_weight    0.2340
Age            10.5000
dtype: float64
```

```
data['Age'] = np.where(data['Age'] < lower_limit, 7, data['Age'])
sns.boxplot(x=data.Age, showfliers=False)
```



7. Check for Categorical columns and perform encoding

In [21]:

```
data.head()
```

| Sex | Length | Diameter | Height | Whole_weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|-----|--------|----------|--------|--------------|----------------|----------------|--------------|------|
| M | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 16.5 |
| M | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 8.5 |
| F | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 10.5 |
| M | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 11.5 |
| I | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 8.5 |

```
from sklearn.preprocessing import LabelEncoder
lab = LabelEncoder()data.Sex = lab.fit_transform(data.Sex)
data.head()
```

```
:
```

| | Sex | Length | Diameter | Height | Whole_weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|------|
| 0 | 2 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 16.5 |
| 1 | 2 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 8.5 |
| 2 | 0 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 10.5 |
| 3 | 2 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 11.5 |
| 4 | 1 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 8.5 |

8. Split the data into dependent and independent variables

```
y = data["Sex"] y.head()
```

In [24]:

```
x=data.drop(columns=["Sex"],axis=1)x.head()
```

| | Length | Diameter | Height | Whole_weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|---|--------|----------|--------|--------------|----------------|----------------|--------------|------|
| 0 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 16.5 |
| 1 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 8.5 |
| 2 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 10.5 |
| 3 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 11.5 |
| 4 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 8.5 |

9. Scale the independent variables

In [25]:

```
from sklearn.preprocessing import scale
X_Scaled = pd.DataFrame(scale(x), columns=x.columns)
X_Scaled.head()
```

| | Length | Diameter | Height | Whole_weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|---|--------|----------|--------|--------------|----------------|----------------|--------------|------|
| 0 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 16.5 |
| 1 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 8.5 |
| 2 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 10.5 |
| 3 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 11.5 |
| 4 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 8.5 |

10. Split the data into training and testing

In [26]:

```
from sklearn.model_selection import train_test_split
```

```
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X_Scaled, y, test_size=0.2,
random_state=0)
```

In [27]:

```
X_Train.shape,X_Test.shape
```

Output

```
((3341, 8), (836, 8))
```

```
Y_Train.shape,Y_Test.shape
```

output

```
((3341,), (836,))
```

In [29]:

```
X_Train.head()
```

| | Length | Diameter | Height | Whole_weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|------|-----------|-----------|-----------|--------------|----------------|----------------|--------------|-----------|
| 3141 | -2.864726 | -2.750043 | -1.423087 | -1.622870 | -1.553902 | -1.583867 | -1.644065 | -1.543234 |
| 3521 | -2.573250 | -2.598876 | -2.020857 | -1.606554 | -1.551650 | -1.565619 | -1.626104 | -1.387181 |
| 883 | 1.132658 | 1.230689 | 0.728888 | 1.145672 | 1.041436 | 0.286552 | 1.538726 | 1.577830 |
| 3627 | 1.590691 | 1.180300 | 1.446213 | 2.164373 | 2.661269 | 2.330326 | 1.377072 | 0.017298 |
| 2106 | 0.591345 | 0.474853 | 0.370226 | 0.432887 | 0.255175 | 0.272866 | 0.906479 | 1.265723 |

X_Test.head()

| | Length | Diameter | Height | Whole_weight | Shucked_weight | Viscera_weight | Shell_weight | Age |
|------|-----------|-----------|-----------|--------------|----------------|----------------|--------------|-----------|
| 3141 | -2.864726 | -2.750043 | -1.423087 | -1.622870 | -1.553902 | -1.583867 | -1.644065 | -1.543234 |
| 3521 | -2.573250 | -2.598876 | -2.020857 | -1.606554 | -1.551650 | -1.565619 | -1.626104 | -1.387181 |
| 883 | 1.132658 | 1.230689 | 0.728888 | 1.145672 | 1.041436 | 0.286552 | 1.538726 | 1.577830 |
| 3627 | 1.590691 | 1.180300 | 1.446213 | 2.164373 | 2.661269 | 2.330326 | 1.377072 | 0.017298 |
| 2106 | 0.591345 | 0.474853 | 0.370226 | 0.432887 | 0.255175 | 0.272866 | 0.906479 | 1.265723 |

Y_Train.head()

```
3141    1
3521    1
883     2
3627    2
2106    2
Name: Sex, dtype: int32
```

```
Y_Test.head()
```

11. Build the Model

```
In [33]:
```

```
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_estimators=10,criterion='entropy')
```

```
In [34]:
```

```
model.fit(X_Train,Y_Train)
```

Output

```
RandomForestClassifier(criterion='entropy', n_estimators=10)
```

12. Train the Model

```
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
```

```
In [38]:
```

```
print('Training accuracy: ',accuracy_score(Y_Train,y_predict_train))
```

Output

```
Training accuracy: 0.9817419934151451
```

13. Test the Model

```
In [39]:
```

```
print('Testing accuracy: ',accuracy_score(Y_Test,y_predict))
```

Output

```
Testing accuracy: 0.5322966507177034
```

14. Measure the performance using Metrics

```
In [40]:
```

```
pd.crosstab(Y_Test,y_predict)
```

```
col_0    0    1    2
```

```
Sex
```

```
0  116   29  104
```

```
1   37  216   38
```

```
2  125   58  113
```

```
print(classification_report(Y_Test,y_predict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.42 | 0.47 | 0.44 | 249 |
| 1 | 0.71 | 0.74 | 0.73 | 291 |
| 2 | 0.44 | 0.38 | 0.41 | 296 |
| accuracy | | | 0.53 | 836 |
| macro avg | 0.52 | 0.53 | 0.53 | 836 |
| weighted avg | 0.53 | 0.53 | 0.53 | 836 |