

## Assignment -2

### Data Visualization & Pre-processing

Assignment Date	22 September 2022
Student Name	Miss Arunadevi A
Student Roll Number	620119106008
Maximum Marks	2 Marks

#### Question-1

1. Dataset downloaded as "model.csv"

2. Load the dataset

```
#importing libraries
import pandas as pd
#load the dataset
df=pd.read_csv("model.csv")
df
```

Out[7]:


	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	96270.64	0
9996	9997	15569692	Johnstone	516	France	Male	35	10	57369.61	1	1	1	101699.77	0
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	42085.58	1
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	92888.52	1
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	38190.78	0

10000 rows x 14 columns

3. Perform Below Visualizations


3.1 Univariate Analysis

```
#scatterplot
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
#load the dataset
df=pd.read_csv("model.csv")
plt.scatter(df.index,df['NumOfProducts'])
plt.show()
```




```
#strip plot
sns.stripplot(y=df['NumOfProducts'])
<AxesSubplot: NumOfProducts>
```

Out[7]



```
#histogram
plt.hist(df['CreditScore'])
```

```
Out[8]:
(array([ 19, 166, 447, 958, 1444, 1866, 1952, 1525, 968,
        655]),
 array([350, 400, 450, 500, 550, 600, 650, 700, 750, 800, 850]),
 <BarContainer object of 10 artists>)
```



```
#boxplot
```


```
sns.boxplot(df['CreditScore'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[10]:
```


```
<AxesSubplot: xlabel='CreditScore'>
```



## 3.2 Bivariate Analysis

In[21]:

```
#scatter plot  
plt.scatter(df.NumOfProducts,df.CreditScore)  
plt.show()
```




In[22]:

```
#strip plot  
sns.stripplot(x=df['NumOfProducts'],y=df['CreditScore'])
```

Out[22]:

```
<AxesSubplot:xlabel='NumOfProducts', ylabel='CreditScore'>
```



## 3.3 Multivariate Analysis

In[12]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
sns.set_style('darkgrid')
sns.set(font_scale=1.3)
```

```
df=pdread_csv('model.csv')
df
```

Out[12]:	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101346.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	6	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737868	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	96270.64	0
9996	9997	15569852	Johnstone	516	France	Male	35	10	57369.61	1	1	1	101699.77	0
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	42065.58	1
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	92888.52	1
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	38190.78	0

10000 rows x 14 columns

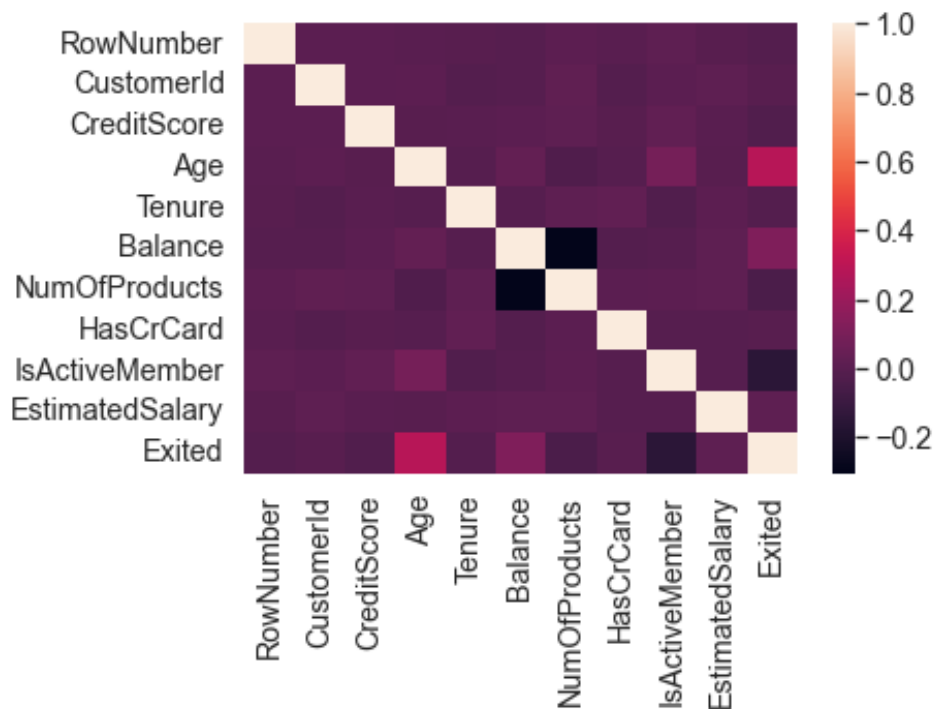
```
#pairplot
sns.pairplot(df);
```



```
sns.heatmap(df.corr())
```

```
<AxeSubplot>
```

Out[14]:



## 4. Perform descriptive statistics on the dataset

*#load the dataset*

`import pandas as pd`

`data=pdread_csv("model.csv")`

`data.head()`

Out[4]:

In [5]:

Out[5]:

In [6]:

data.tail()

Out[6]:

In [7]:

data.tail(10)

Out[7]:

In [8]:

data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column              Non-Null Count  Dtype  
---  --
 0   RowNumber           10000 non-null  int64  
 1   CustomerId          10000 non-null  int64  
 2   Surname             10000 non-null  object  
 3   CreditScore         10000 non-null  int64  
 4   Geography           10000 non-null  object  
 5   Gender              10000 non-null  object  
 6   Age                 10000 non-null  int64  
 7   Tenure              10000 non-null  int64  
 8   Balance             10000 non-null  float64 
 9   NumOfProducts       10000 non-null  int64  
10   HasCrCard           10000 non-null  int64  
11   IsActiveMember      10000 non-null  int64  
12   EstimatedSalary     10000 non-null  float64 
13   Exited              10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

In [9]: data.shape
Out[9]: (10000, 14)

In [10]: data.head()
Out[10]:
RowNumber    5.000000e+03
CustomerId   1.565070e+07
CreditScore   6.595280e+02
Age           3.892380e+01
Tenure        5.012800e+00
Balance       7.446350e+04
NumOfProducts 1.530200e+00
HasCrCard     7.055000e-01
IsActiveMember 5.350000e-01
```

```
C:\Users\jamar vijay\AppData\Local\Temp\ipykernel_5088\51206388.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError.  Select only valid columns before calling the reduction.
data.head()
Out[11]:
RowNumber    5.000000e+03
CustomerId   1.565070e+07
CreditScore   6.520000e+02
Age           3.700000e+01
Tenure        8.000000e+00
Balance       9.710140e+04
NumOfProducts 1.000000e+00
HasCrCard     1.000000e+00
IsActiveMember 1.000000e+00
EstimatedSalary 2.002230e+05
Exited        0.000000e+00
dtype: float64

In [12]: data.head()
Out[12]:
RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  EstimatedSalary  Exited
0          1  1565701    Smith      850.0    France  Male   37.0   2.0   0.0           1.0           1.0           1.0      245124.92   0.0
1          2  1565706    NaN           NaN           NaN  NaN    NaN    NaN    NaN           NaN           NaN           NaN           NaN
2          3  1565714    NaN           NaN           NaN  NaN    NaN    NaN    NaN           NaN           NaN           NaN           NaN
3          4  1565779    NaN           NaN           NaN  NaN    NaN    NaN    NaN           NaN           NaN           NaN           NaN
4          5  1565790    NaN           NaN           NaN  NaN    NaN    NaN    NaN           NaN           NaN           NaN           NaN
...         ...         ...         ...         ...         ...  ...   ...   ...   ...           ...           ...           ...           ...
9995       9995  15019020  NaN           NaN           NaN  NaN    NaN    NaN    NaN           NaN           NaN           NaN           NaN
9996       9997  15019025  NaN           NaN           NaN  NaN    NaN    NaN    NaN           NaN           NaN           NaN           NaN
9997       9998  15019050  NaN           NaN           NaN  NaN    NaN    NaN    NaN           NaN           NaN           NaN           NaN
9998       9999  15019052  NaN           NaN           NaN  NaN    NaN    NaN    NaN           NaN           NaN           NaN           NaN
9999      10000  15019060  NaN           NaN           NaN  NaN    NaN    NaN    NaN           NaN           NaN           NaN           NaN

10000 rows x 14 columns
```

we can see the wealthier passengers in the higher classes tend to be older, which makes sense average age

values to impute based on Pclass for Age

```
In [29]: def impute_age(cols):
        Age=cols[0]
        Pclass=cols[1]

        if pd.isnull(Age):

            if Pclass==1:
                return 37

            elif Pclass ==2:
                return 29

            else:
                return 24
        else:
            return Age
```

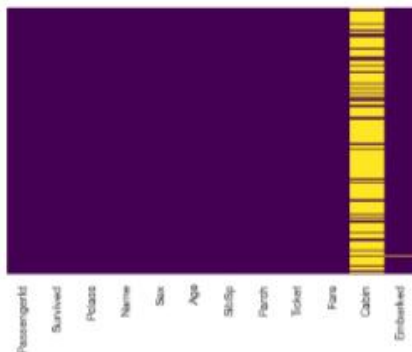
Now Apply This Function

```
In [30]: train['Age'] = train[['Age','Pclass']].apply(impute_age,axis=1)
```

Now let's check Heatmap Again..

```
In [31]: sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

Out[31]: <AxesSubplot:>



Now The Age Missing Values Can be Handled.

## 6. Find the outliers and replace the outliers

```
In [16]: #import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

#load the dataset
df=pd.read_csv('model.csv')
df
```


```
Out[16]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	
...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	

10000 rows x 14 columns




Out[18]: <AxesSubplot:>




In [19]: `sns.countplot(x='SibSp', data=train)`

Out[19]: <AxesSubplot: xlabel='SibSp', ylabel='count'>




In [20]: `train['Fare'].hist(color='green', bins=40, figsize=(8,4))`

Out[20]: <AxesSubplot:>




In [21]: `#Data cleaning  
plt.figure(figsize=(12,7))  
sns.boxplot(x='Pclass', y='Age', data=train, palette='winter')`

Out[21]: <AxesSubplot: xlabel='Pclass', ylabel='Age'>




```
In [11]: #plotting outliers
sns.boxplot(df["CreditScore"])

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
Out[11]: <AxesSubplot: xlabel='CreditScore'>
```



```
In [39]: sns.boxplot(df["Age"])

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
Out[39]: <AxesSubplot: xlabel='Age'>
```



```
In [12]: qnt=df.quantile(q=[0.75,0.25])
qnt
```

```
Out[12]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0.75	7500.25	15753233.75	718.0	44.0	7.0	127644.24	2.0	1.0	1.0	149388.2475
0.25	2500.75	15628928.25	584.0	32.0	3.0	0.00	1.0	0.0	0.0	91002.1100

```
In [14]: iqr = qnt.loc[0.75]-qnt.loc[0.25] #iqr calculations
iqr
```

```
Out[14]:
```

RowNumber	4999.5000
CustomerId	124705.5000
CreditScore	134.0000
Age	12.0000
Tenure	4.0000
Balance	127644.2400
NumOfProducts	1.0000
HasCrCard	1.0000
IsActiveMember	1.0000
EstimatedSalary	98386.1375
Exited	0.0000
dtype:	float64

```
In [26]: #lower extreme values
lower=qnt.loc[0.25] - 1.5*iqr
lower
```

```
Out[26]:
```

RowNumber	-4.998500e+03
CustomerId	1.544147e+07
CreditScore	3.830000e+02
Age	1.400000e+01
Tenure	-3.000000e+00
Balance	-1.914664e+05
NumOfProducts	-5.000000e-01
HasCrCard	-1.500000e+00
IsActiveMember	-1.500000e+00
EstimatedSalary	-9.657710e+04
Exited	0.000000e+00
dtype:	float64

```
In [27]: #upper extreme values
upper=qnt.loc[0.75] + 1.5*iqr

upper
```

```
Out[27]: RowNumber      1.499950e+04
CustomerId    1.594029e+07
CreditScore   9.190000e+02
Age           6.200000e+01
Tenure        1.300000e+01
Balance       3.191106e+05
NumOfProducts 3.500000e+00
HasCrCard     2.500000e+00
IsActiveMember 2.500000e+00
EstimatedSalary 2.969675e+05
Exited        0.000000e+00
dtype: float64
```

```
In [18]: df.mean()
```

C:\Users\janar vijay\AppData\Local\Temp\ipykernel\_10016\3698961737.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
df.mean()
Out[18]: RowNumber      5.000500e+03
CustomerId    1.569094e+07
CreditScore   6.505288e+02
Age           3.892180e+01
Tenure        5.012800e+00
Balance       7.648589e+04
NumOfProducts 1.530200e+00
HasCrCard     7.055000e-01
IsActiveMember 5.151000e-01
EstimatedSalary 1.000902e+05
Exited        2.037000e-01
dtype: float64
```

## Replacing outlier


```
In [45]: #import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

#load the dataset
df=pd.read_csv('model.csv')
df['CreditScore']=np.where(df['CreditScore']<400,402,df['CreditScore'])
```

```
In [49]: #remove outlier on the CreditScore column
sns.boxplot(df["CreditScore"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
<AxesSubplot:xlabel='CreditScore'>
```




```
In [50]: #remove outlier on the Age column
df['Age']=np.where(df['Age']>60,50,df['Age'])
```

```
In [51]: sns.boxplot(df["Age"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
<AxesSubplot:xlabel='Age'>
```



## 7. Check for Categorical columns and perform encoding

```
In [53]: df.head()
```

```
Out[53]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15674602	Hargrave	619	France	Female	42	2	0.00	1	1	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1

### encoding

```
In [57]: #manually handling categorincal data
dff['Gender'].replace({'Female':1,'Male':2},inplace=True)
dff.head()
```

```
Out[57]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15674602	Hargrave	619	France	1	42	2	0.00	1	1	1
1	2	15647311	Hill	608	Spain	1	41	1	83807.86	1	0	0
2	3	15619304	Onio	502	France	1	42	8	159660.80	3	1	1
3	4	15701354	Boni	699	France	1	39	1	0.00	2	0	0
4	5	15737888	Mitchell	850	Spain	1	43	2	125510.82	1	1	1

```
In [62]: dff['Geography'].replace({'France':100,'Spain':200},inplace=True)
dff.head()
```

```
Out[62]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15674602	Hargrave	619	100	1	42	2	0.00	1	1	1
1	2	15647311	Hill	608	200	1	41	1	83807.86	1	0	0
2	3	15619304	Onio	502	100	1	42	8	159660.80	3	1	1
3	4	15701354	Boni	699	100	1	39	1	0.00	2	0	0
4	5	15737888	Mitchell	850	200	1	43	2	125510.82	1	1	1

```
In [60]: #dummy variable function
df_main=pd.get_dummies(df,columns=['Geography'])
df_main
```

```
Out[60]:
```

	RowNumber	CustomerId	Surname	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15674602	Hargrave	619	1	42	2	0.00	1	1	1
1	2	15647311	Hill	608	1	41	1	83807.86	1	0	0
2	3	15619304	Onio	502	1	42	8	159660.80	3	1	1
3	4	15701354	Boni	699	1	39	1	0.00	2	0	0
4	5	15737888	Mitchell	850	1	43	2	125510.82	1	1	1
...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	731	2	39	5	0.00	2	1	0
9996	9997	15569892	Johnstone	516	2	35	10	57369.61	1	1	1
9997	9998	15584532	Liu	709	1	36	7	0.00	1	0	1
9998	9999	15682355	Sabbatini	732	2	42	3	75075.31	2	1	0
9999	10000	15628319	Walker	792	1	28	4	130142.79	1	1	0

10000 rows x 16 columns

## 8. Split the data into dependent and independent variables.

```
In [64]: #target variable or dependent variable.
import pandas as pd
df=pd.read_csv('model.csv')
y=df['EstimatedSalary']
y.head()
```

```
Out[64]:
```

0 101348.88  
1 112542.58  
2 113931.57  
3 93826.63  
4 79084.10  
Name: EstimatedSalary, dtype: float64

```
In [74]: #independent variables
x = df.drop(columns=['EstimatedSalary'],axis=1)
x.head()
```

```
Out[74]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15674602	Hargrave	619	France	Female	42	2	0.00	1	1	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1



